

EXPERIMENT NO: 9

Social Network Ads Classification using Logistic Regression

Aim:

To build a Logistic Regression model that predicts whether a user purchases a product based on their age and estimated salary.

Algorithm:

- 1. Load Data: Read the Social Network Ads dataset using pandas.**
- 2. Select Features and Labels:**
- 3. Features → Age and Estimated Salary**
- 4. Label → Purchased (0 or 1)**
- 5. Model Optimization:**
- 6. Run multiple train-test splits with different random states (1–400).**
- 7. Identify the state where the test accuracy exceeds the training accuracy.**
- 8. Train Final Model:**
- 9. Split data with the best random state (306).**
- 10. Train the Logistic Regression model.**
- 11. Evaluate Model:**
- 12. Display training and testing accuracy.**
- 13. Generate classification report to assess precision, recall, and F1-score.**

Program:

```
[1]: import pandas as pd
import numpy as np
df=pd.read_csv("C:/Users/vijay/Downloads/Social_Network_Ads.csv")
df
```

```
[1]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
[2]: df.head()
```

```
[2]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
[3]: features=df.iloc[:,[2,3]].values
      label=df.iloc[:,4].values
      features
```

```
[3]: array([[ 19, 19000],
             [ 35, 20000],
             [ 26, 43000],
             [ 27, 57000],
             [ 19, 76000],
             [ 27, 58000],
             [ 27, 84000],
             [ 32, 150000],
             [ 25, 33000],
             [ 35, 65000],
             [ 26, 80000],
             [ 26, 52000],
             [ 20, 86000],
             [ 32, 18000],
             [ 18, 82000],
             [ 29, 80000],
             [ 47, 25000],
             [ 45, 26000])
```

```
[4]: label
```

```
[4]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
          0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,
          1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
          1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1,
          0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1,
          1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,
          0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
          1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1,
          0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,
          1, 1, 0, 1])
```

```
[5]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
```

```
[8]: for i in range(1,401):
      xtrain,xtest,ytrain,ytest=train_test_split(features,label,test_size=0.2,random_state=i)
      model=LogisticRegression()
      model.fit(xtrain,ytrain)
      train_score=model.score(xtrain,ytrain)
      test_score=model.score(xtest,ytest)
      if test_score>train_score:
          print("Test: {:.3f}, Train: {:.3f}, Random State: {}".format(test_score,train_score,i))
```

```
Test: 0.900, Train: 0.841, Random State: 10
Test: 0.863, Train: 0.856, Random State: 14
Test: 0.850, Train: 0.844, Random State: 15
Test: 0.863, Train: 0.856, Random State: 16
Test: 0.875, Train: 0.834, Random State: 18
Test: 0.850, Train: 0.844, Random State: 19
Test: 0.875, Train: 0.844, Random State: 20
Test: 0.863, Train: 0.834, Random State: 21
Test: 0.875, Train: 0.841, Random State: 22
Test: 0.875, Train: 0.841, Random State: 24
Test: 0.850, Train: 0.834, Random State: 26
Test: 0.850, Train: 0.841, Random State: 27
Test: 0.863, Train: 0.834, Random State: 30
```

```
[9]: xtrain,xtest,ytrain,ytest=train_test_split(features,label,test_size=0.2,random_state=306)
      fmodel=LogisticRegression()
      fmodel.fit(xtrain,ytrain)
```

```
[9]: ▾ LogisticRegression ⓘ ⓘ
      LogisticRegression()
```

```
[10]: print(fmodel.score(xtrain,ytrain))
      print(fmodel.score(xtest,ytest))

      0.8375
      0.9125
```

```
[12]: from sklearn.metrics import classification_report
      print(classification_report(label,fmodel.predict(features)))
```

	precision	recall	f1-score	support
0	0.86	0.91	0.89	257
1	0.83	0.74	0.78	143
accuracy			0.85	400
macro avg	0.85	0.83	0.84	400
weighted avg	0.85	0.85	0.85	400

Result:

The Logistic Regression model effectively predicts whether a customer will purchase based on age and salary. The model shows balanced training and testing accuracy, and the classification report confirms good overall performance.